




Trustworthy Traffic Signal Control with Deep Reinforcement Learning

Cheuk Hong Tang¹ Jana Speldrich¹ Dunkun Wei¹


¹Department of Computer Science, National Yang Ming Chiao Tung University

Abstract

Traffic signal control in urban environments requires balancing vehicle throughput and pedestrian safety. To address this challenge, we propose an **Action-Constrained Proximal Policy Optimisation (PPO)** framework that integrates pedestrian safety requirements through feasible-action masking during policy execution. The proposed method is compared with five widely used traffic-control baselines — Fixed-Time, Max Pressure, SOTL, DQN-AR, and SPRe+ — in both a single-intersection task and a 3×3 grid network simulated in SUMO. Through these experiments, we examine the potential of safe and explainable deep reinforcement learning for trustworthy traffic signal control in smart-city environments.



(a) 1x1 env




(b) 3x3 env

Research Objectives

- Objective 1:** Design an action-constrained RL agent that enforces pedestrian safety constraints *by construction*, guaranteeing zero violations without relying on reward shaping alone.
- Objective 2:** Benchmark the proposed agent against rule-based baselines (Fixed-Time, Max Pressure, SOTL) and action-constrained learned agents (DQN-AR, SPRe+) on efficiency and safety metrics across two network scales.
- Objective 3:** Validate the computational advantage of acceptance-rejection masking over QP-projection (SPRe+) in growing action spaces.

Methodology

The agent operates within an **Action-Constrained MDP**. At every step, the constraint checker computes the feasible action set $\mathcal{C}(s)$ by removing phases that violate any hard constraint. PPO selects actions via a **masked softmax** restricted to $\mathcal{C}(s)$, so infeasible actions receive zero probability and can never be executed.



Experimental Setup

Simulation Scenarios

Parameter	1×1 Scenario	3×3 Scenario
Network topology	Single 4-way intersection	3×3 grid
Controlled TLS	1	B1 (center node)
Simulation length	3 600 s	3 600 s
Training episodes	100	100
Evaluation episodes	5	5
State dimension	21	31
Action space	2 phases	4 phases

Agents & Baselines

- Fixed-Time:** Fixed 30 s NS/EW cycle. No traffic awareness.
- Max Pressure:** Selects phase maximising queue pressure. Vehicle-centric; no pedestrian safety.
- SOTL:** Switches when red-side queue $\Phi_{\text{red}}(t) \geq \kappa=5$ and $\tau \geq \tau_{\text{min}}$.
- DQN-AR:** DQN with acceptance-rejection safety filter. Zero violations; $\mathcal{O}(1)$ per-step cost.
- SPRe+ :** Projects policy onto feasible simplex via QP. $\mathcal{O}(|\mathcal{A}|^2)$; equivalent to DQN-AR here.
- Action-Constrained PPO (proposed):** Masked softmax over $\mathcal{C}(s)$. Learns safe behaviour during training.

Results & Discussion

Table 1 — 1×1 Single Intersection (5 episodes × 3 600 s)

Agent	Avg Wait (s)	Avg Queue	Viol. Rate	Reward
Action-Constr. PPO	60.28	4.27	0.0%	−217,026
DQN-AR	65.85	4.47	0.0%	−237,073
SPRe+ (DQN policy)	65.85	4.47	0.0%	−237,073
SOTL ($\kappa=5$)	71.64	4.78	18.6%	−291,292
Max Pressure [†]	55.63	4.04	53.5%	−297,984
Fixed-Time (30 s)	107.58	5.61	57.2%	−490,231

[†]Lowest raw wait but 53.5% violation rate — unsafe for real deployment.

Table 2 — 3×3 Grid, Intersection B1 (5 episodes × 3 600 s)

Agent	Avg Wait (s)	Avg Queue	Viol. Rate	Reward
Action-Constr. PPO	1,291.4	78.67	0.0%	−4,649,048
DQN-AR	1,350.6	79.61	0.0%	−4,862,170
SPRe+ (DQN policy)	1,350.6	79.61	0.0%	−4,862,170
Max Pressure	1,380.1	79.58	55.8%	−5,068,558
SOTL ($\kappa=5$)	1,406.0	80.65	71.0%	−5,189,343
Fixed-Time (30 s)	1,604.2	80.27	69.1%	−5,899,362

Key Findings

- High Efficiency:** PPO reduces vehicle wait times by **8.46%** (1×1) and **4.38%** (3×3) compared to DQN-AR.
- Guaranteed Safety:** Learned agents achieve **0% violations**. In contrast, rule-based heuristics violate safety rules in **19–71%** of steps (e.g., Max Pressure at 53.5%).
- Network Scalability:** In the partially-controlled 3×3 grid, PPO maintains a strong lead over Fixed-Time, narrowing the gap to Max Pressure (−6.4%).

Conclusions & Key Contributions

- Addressing Benchmark Limitations:** While traditional heuristics like *Max Pressure* and *SOTL* are vehicle-centric and lack pedestrian safety, and state-of-the-art *SPRe+* guarantees zero violations but suffers from high computation cost (QP projection overhead), our approach bridges these gaps.
- Optimal Safety-Efficiency Trade-off:** The proposed **Action-constrained PPO** achieves the lowest waiting time among safe agents with zero pedestrian-safety violations guaranteed by construction.
- Eliminated Computational Overhead:** By enforcing feasible-action masking during training, our method completely eliminates the post-hoc projection overhead of *SPRe+* and the resampling loops of *DQN-AR*, ensuring scalability for larger networks.
- Scalability & Future Horizons:** Longer training is expected to further widen the performance gap in complex 3 × 3 grids. Future steps include full 9-intersection multi-agent control and evaluation on real-world empirical data.

References

- [1] P. Varaiya, “Max pressure control of a network of signalized intersections,” *Transportation Research Part C*, vol. 36, pp. 177–195, 2013.
- [2] C. Gershenson, “Self-organizing traffic lights,” *Complex Systems*, vol. 16, no. 1, pp. 29–53, 2005.
- [3] Y.-H. Hung et al., “SPRe+: Safe policy with repair for action-constrained reinforcement learning,” 2025.